

# An introduction to Apache Avro

Douglas Creager  
RedJack, LLC

Mil-OSS WG3  
August 30, 2011

# What is it?

---

A framework for:

- *Modeling* your data
- *Serializing* your data
- Making *remote procedure calls*

# What is it?

---

A framework for:

- *Modeling* your data
- *Serializing* your data
- Making *remote procedure calls*

# Benefits

---

- Intuitive modeling paradigm
- Efficient binary encoding
- Forwards- and backwards-compatibility via *schema resolution*

# Modeling paradigms

# Modeling paradigms

## Relational model

PEOPLE			PARENTS	
person_id	name	dob	parent	child
1	Buford	11/19/1953	1	3
2	Wilhelmina	04/27/1956	2	3
3	Sanjay	06/12/1980	1	4
4	Jill	03/08/1982	2	4

# Modeling paradigms

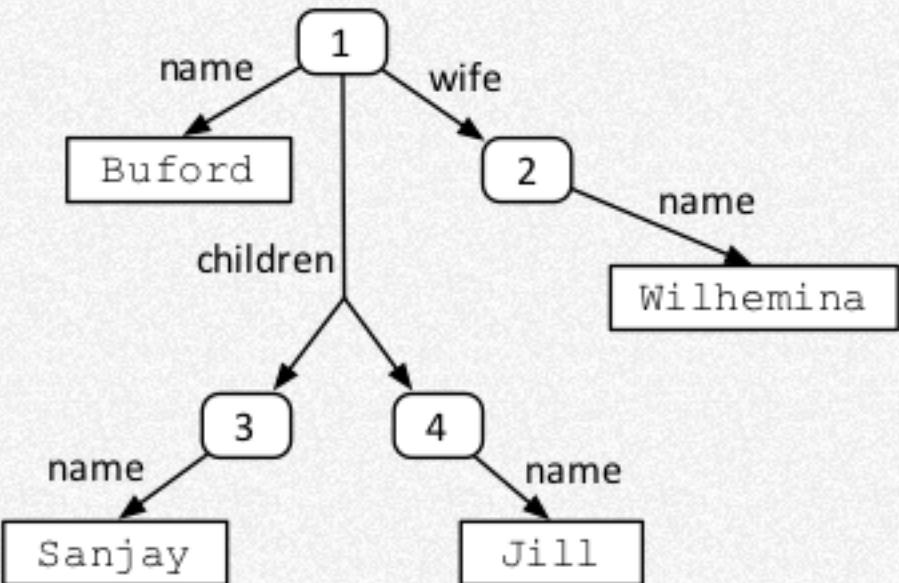
## XML – Extensible Markup Language

```
<?xml version="1.0">
<people>
  <person id="1">
    <name>Buford</name>
    <wife>
      <person id="2">
        <name>Wilhelmina</name>
      </person>
    </wife>
    <children>
      <person id="3">
        <name>Sanjay</name>
      </person>
      <person id="3">
        <name>Jill</name>
      </person>
    </children>
  </person>
</people>
```

# Modeling paradigms

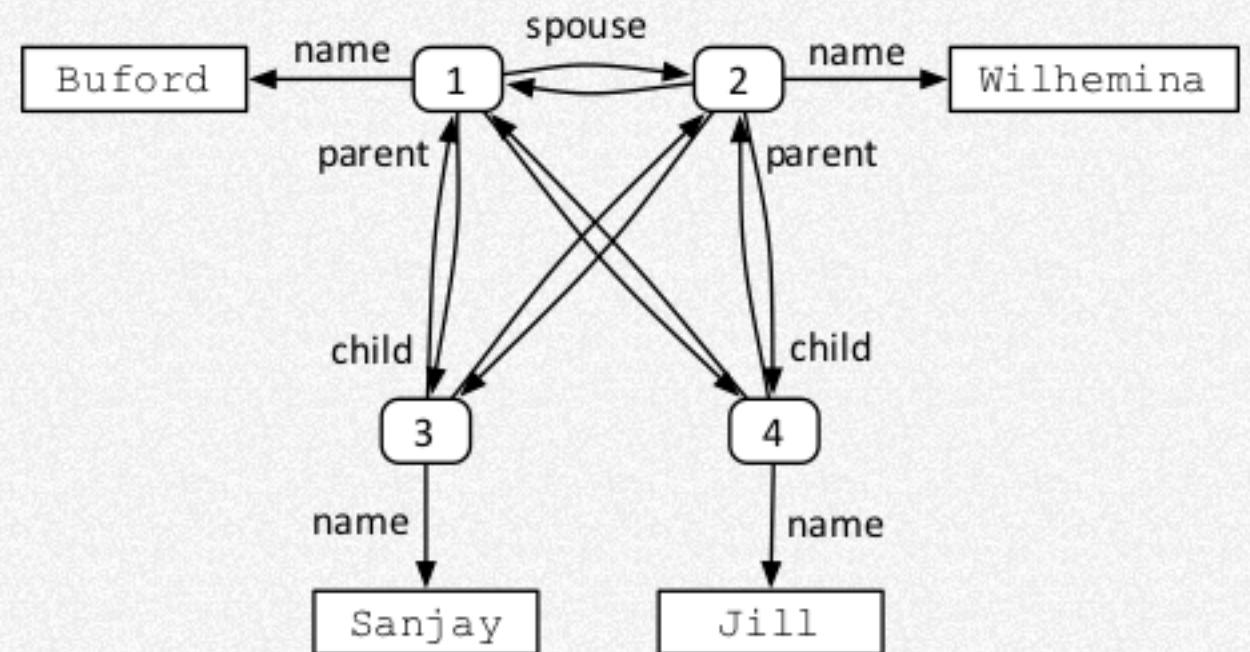
## XML – Extensible Markup Language

```
<?xml version="1.0">
<people>
  <person id="1">
    <name>Buford</name>
    <wife>
      <person id="2">
        <name>Wilhelmina</name>
      </person>
    </wife>
    <children>
      <person id="3">
        <name>Sanjay</name>
      </person>
      <person id="3">
        <name>Jill</name>
      </person>
    </children>
  </person>
</people>
```



# Modeling paradigms

## Semantic Web



# Impedence mismatch!

---

What data structures do we use in our programs?

# Impedence mismatch!

---

What data structures do we use in our programs?

- *scalars*
  - integers (0, 12, -1, 65,536)
  - floats (0.0, 43.12, 1.234e-17)
  - strings ("Buford")
  - booleans (true, false)
  - null

# Impedence mismatch!

---

What data structures do we use in our programs?

- *scalars*
  - integers (0, 12, -1, 65,536)
  - floats (0.0, 43.12, 1.234e-17)
  - strings ("Buford")
  - booleans (true, false)
  - null
- *arrays* (lists, sequences, vectors)

# Impedence mismatch!

---

What data structures do we use in our programs?

- *scalars*
  - integers (0, 12, -1, 65,536)
  - floats (0.0, 43.12, 1.234e-17)
  - strings ("Buford")
  - booleans (true, false)
  - null
- *arrays* (lists, sequences, vectors)
- *maps* (dictionaries, hashes, associative arrays)

# Impedence mismatch!

---

What data structures do we use in our programs?

- *scalars*
  - integers (0, 12, -1, 65,536)
  - floats (0.0, 43.12, 1.234e-17)
  - strings ("Buford")
  - booleans (true, false)
  - null
- *arrays* (lists, sequences, vectors)
- *maps* (dictionaries, hashes, associative arrays)
- *records* (structs, objects, tuples)

# Modeling paradigms

## JSON – JavaScript Object Notation

```
{  
  "id": 1,  
  "name": "Buford",  
  "wife": {  
    "id": 2,  
    "name": "Wilhelmina"  
  },  
  "children": [  
    {  
      "id": 3,  
      "name": "Sanjay"  
    },  
    {  
      "id": 4,  
      "name": "Jill"  
    }  
  ]  
}
```

- Uses same data structures as our programs

# Serialization and compatibility

# Serialization

---

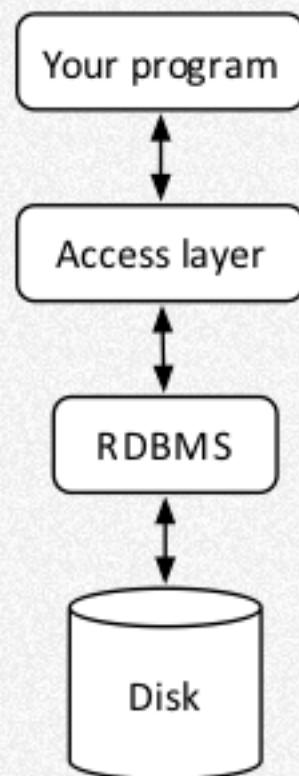
## Custom binary formats

- Incredibly efficient
- Even more incredibly difficult to maintain
  - Cumbersome to write
  - Keeping different implementations in sync
  - Forwards and backwards compatibility? Feh.

# Serialization

## Relational databases

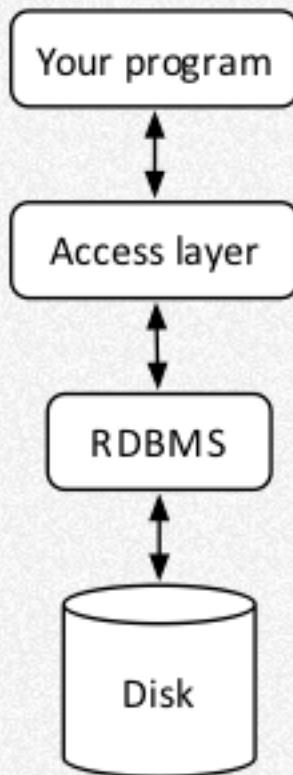
- A great solution:  
*You don't care!*



# Serialization

## Relational databases

- A great solution:  
*You don't care!*
- But...  
you lose interoperability



# Serialization

## XML to the rescue!

```
<?xml version="1.0">
<people>
  <person id="1">
    <name>Buford</name>
    <wife>
      <person id="2">
        <name>Wilhelmina</name>
        </person>
      </wife>
      <children>
        <person id="3">
          <name>Sanjay</name>
        </person>
        <person id="3">
          <name>Jill</name>
        </person>
      </children>
    </person>
  </people>
```

- Human-readable, standardized syntax
- Let libxml/expat/whatever handle the parsing

# Serialization

## XML to the rescue!

```
<?xml version="1.0">
<people>
  <person id="1">
    <name>Buford</name>
    <age>57</age>
    <wife>
      <person id="2">
        <name>Wilhelmina</name>
        <age>55</age>
      </person>
    </wife>
    <children>
      <person id="3">
        <name>Sanjay</name>
        <age>31</age>
      </person>
      <person id="3">
        <name>Jill</name>
        <age>29</age>
      </person>
    </children>
  </person>
</people>
```

- Human-readable, standardized syntax
- Let libxml/expat/whatever handle the parsing
- Compatibility:  
*New tags don't break the parsing*

# Serialization

## “Human-readable”?

```
<?xml version="1.0"?>

<!-- root element wsdl:definitions defines set of related services -->
<wsdl:definitions name="EndorsementSearch"
    targetNamespace="http://namespaces.snowboard-info.com"
    xmlns:es="http://www.snowboard-info.com/EndorsementSearch.wsdl"
    xmlns:esxsd="http://schemas.snowboard-info.com/EndorsementSearch.xsd"
    xmlns:soap="http://schemas.xmlsoap.org/wsdl/soap/"
    xmlns:wsdl="http://schemas.xmlsoap.org/wsdl/">

    <!-- wsdl:types encapsulates schema definitions of communication types; here
using xsd -->
    <wsdl:types>

        <!-- all type declarations are in a chunk of xsd -->
        <xsd:schema targetNamespace="http://namespaces.snowboard-info.com"
            xmlns:xsd="http://www.w3.org/1999/XMLSchema">

            <!-- xsd definition: GetEndorsingBoarder [manufacturer string, model
string] -->
            <xsd:element name="GetEndorsingBoarder">
                <xsd:complexType>
                    <xsd:sequence>
                        <xsd:element name="manufacturer" type="string"/>
                        <xsd:element name="model" type="string"/>
                    </xsd:sequence>
                </xsd:complexType>
            </xsd:element>
        </xsd:schema>
    </wsdl:types>
</wsdl:definitions>
```

From <http://www.w3.org/2001/03/14-annotated-WSDL-examples.html>

# Serialization

## “Human-readable”?

```
<!-- xsd definition: GetEndorsingBoarderResponse [ ... endorsingBoarder
string ... ] -->
<xsd:element name="GetEndorsingBoarderResponse">
  <xsd:complexType>
    <xsd:all>
      <xsd:element name="endorsingBoarder" type="string"/>
    </xsd:all>
  </xsd:complexType>
</xsd:element>

<!-- xsd definition: GetEndorsingBoarderFault [ ... errorMessage string ... ]
-->
<xsd:element name="GetEndorsingBoarderFault">
  <xsd:complexType>
    <xsd:all>
      <xsd:element name="errorMessage" type="string"/>
    </xsd:all>
  </xsd:complexType>
</xsd:element>

</xsd:schema>
</wsdl:types>

<!-- wsdl:message elements describe potential transactions -->

<!-- request GetEndorsingBoarderRequest is of type GetEndorsingBoarder -->
<wsdl:message name="GetEndorsingBoarderRequest">
  <wsdl:part name="body" element="esxsd:GetEndorsingBoarder"/>
</wsdl:message>
```

From <http://www.w3.org/2001/03/14-annotated-WSDL-examples.html>

# Serialization

## “Human-readable”?

```
<!-- response GetEndorsingBoarderResponse is of type  
GetEndorsingBoarderResponse -->  
<wsdl:message name="GetEndorsingBoarderResponse">  
  <wsdl:part name="body" element="esxsd:GetEndorsingBoarderResponse"/>  
</wsdl:message>  
  
<!-- wsdl:portType describes messages in an operation -->  
<wsdl:portType name="GetEndorsingBoarderPortType">  
  
  <!-- the value of wsdl:operation eludes me -->  
  <wsdl:operation name="GetEndorsingBoarder">  
    <wsdl:input message="es:GetEndorsingBoarderRequest"/>  
    <wsdl:output message="es:GetEndorsingBoarderResponse"/>  
    <wsdl:fault message="es:GetEndorsingBoarderFault"/>  
  </wsdl:operation>  
</wsdl:portType>  
  
<!-- wsdl:binding states a serialization protocol for this service -->  
<wsdl:binding name="EndorsementSearchSoapBinding"  
  type="es:GetEndorsingBoarderPortType">  
  
  <!-- leverage off soap:binding document style @@@(no wsdl:foo pointing at the  
  soap binding) -->  
  <soap:binding style="document"  
    transport="http://schemas.xmlsoap.org/soap/http"/>  
  
  <!-- semi-opaque container of network transport details classed by  
  soap:binding above @@@ -->  
  <wsdl:operation name="GetEndorsingBoarder">
```

From <http://www.w3.org/2001/03/14-annotated-WSDL-examples.html>

# Serialization

## “Human-readable”?

```
<!-- again bind to SOAP? @@@ -->
<soap:operation soapAction="http://www.snowboard-
info.com/EndorsementSearch"/>

    <!-- furthur specify that the messages in the wsdl:operation
"GetEndorsingBoarder" use SOAP? @@@ -->
    <wsdl:input>
        <soap:body use="literal"
            namespace="http://schemas.snowboard-
info.com/EndorsementSearch.xsd"/>
    </wsdl:input>
    <wsdl:output>
        <soap:body use="literal"
            namespace="http://schemas.snowboard-
info.com/EndorsementSearch.xsd"/>
    </wsdl:output>
    <wsdl:fault>
        <soap:body use="literal"
            namespace="http://schemas.snowboard-
info.com/EndorsementSearch.xsd"/>
    </wsdl:fault>
    </wsdl:operation>
</wsdl:binding>

<!-- wsdl:service names a new service "EndorsementSearchService" -->
<wsdl:service name="EndorsementSearchService">
    <wsdl:documentation>snowboarding-info.com Endorsement
Service</wsdl:documentation>
```

From <http://www.w3.org/2001/03/14-annotated-WSDL-examples.html>

# Serialization

## “Human-readable”?

```
<!-- connect it to the binding "EndorsementSearchSoapBinding" above -->
<wsdl:port name="GetEndorsingBoarderPort"
             binding="es:EndorsementSearchSoapBinding">

    <!-- give the binding an network address -->
    <soap:address location="http://www.snowboard-info.com/EndorsementSearch"/>
</wsdl:port>
</wsdl:service>

</wsdl:definitions>
```

From <http://www.w3.org/2001/03/14-annotated-WSDL-examples.html>

# Serialization

## “Human-readable”?

The equivalent in C:

```
const char *
GetEndorsingBoarder(const char *manufacturer,
                     const char *model);
```

# Serialization

## “Human-readable”?

The equivalent in C:

```
const char *
GetEndorsingBoarder(const char *manufacturer,
                     const char *model);
```

- Less and less human-readable with complex types
- *So much* overhead...

# Apache Avro (finally)

# Data modeling

```
{  
  "type": "record",  
  "name": "person",  
  "fields": [  
    { "name": "name",  
      "type": "string"},  
    { "name": "age",  
      "type": "int"},  
    { "name": "spouse",  
      "type": [ "null", "person" ]},  
    { "name": "children",  
      "type": { "type": "array",  
                "items": "person" } }  
  ]  
}
```

# Data modeling

```
{  
    "type": "record",  
    "name": "person",  
    "fields": [  
        { "name": "name",  
            "type": "string"},  
        { "name": "age",  
            "type": "int"},  
        { "name": "spouse",  
            "type": [ "null", "person" ]},  
        { "name": "children",  
            "type": { "type": "array",  
                    "items": "person" } }  
    ]  
}
```

- Schemas are encoded in JSON

# Data modeling

```
{  
  "type": "record",  
  "name": "person",  
  "fields": [  
    {"name": "name",  
     "type": "string"},  
    {"name": "age",  
     "type": "int"},  
    {"name": "spouse",  
     "type": ["null", "person"]},  
    {"name": "children",  
     "type": {"type": "array",  
              "items": "person"}  
  ]  
}
```

- Schemas are encoded in JSON
- Schema language uses standard programming data structures

# Data modeling

```
{  
  "type": "record",  
  "name": "person",  
  "fields": [  
    {"name": "name",  
     "type": "string"},  
    {"name": "age",  
     "type": "int"},  
    {"name": "spouse",  
     "type": ["null", "person"]},  
    {"name": "children",  
     "type": {"type": "array",  
              "items": "person"}  
  ]  
}
```

- Schemas are encoded in JSON
- Schema language uses standard programming data structures
- Including:
  - scalars:
    - boolean, double, float, int, long, null, string
    - enums
    - variable-size byte arrays
    - fixed-size byte arrays
  - arrays and maps
  - records
  - unions

# Serialization

---

Key assumption:

*The writer schema will always be available when you read the data.*

# Serialization

---

## Scalars

- int and long: “varint” encoding
  - 0 → 00
  - 1 → 01
  - 1 → 02
  - 64 → 80 01
- double and float: raw 4- or 8-byte IEEE 754 encoding
- boolean: 00 or 01
- null: zero bytes
- bytes and string:  
varint-encoded length prefix, followed by binary data
- enum: encoded as an integer, not a string
- fixed: no length prefix needed

# Serialization

---

## Compound values

# Serialization

---

## Compound values

- array and map
  - element count followed by elements
  - “blocks” allow you to serialize without knowing full array length in advance

# Serialization

---

## Compound values

- array and map
  - element count followed by elements
  - “blocks” allow you to serialize without knowing full array length in advance
- union
  - integer discriminant followed by branch value

# Serialization

---

## Compound values

- array and map
  - element count followed by elements
  - “blocks” allow you to serialize without knowing full array length in advance
- union
  - integer discriminant followed by branch value
- record
  - field values in order
  - *no field names or indices*
  - no overhead for records

# Serialization

```
{  
  "name": "Buford",  
  "age": 57,  
  "wife": {  
    "person": {  
      "name": "Wilhelmina",  
      "age": 55  
    }  
  },  
  "children": [  
    {  
      "name": "Sanjay",  
      "age": 31  
    },  
    {  
      "name": "Jill",  
      "age": 29  
    }  
  ]  
}
```

There's also a JSON encoding for  
human-readability

# Serialization

```
{  
  "name": "Buford",  
  "age": 57,  
  "wife": {  
    "person": {  
      "name": "Wilhelmina",  
      "age": 55  
    }  
  },  
  "children": [  
    {  
      "name": "Sanjay",  
      "age": 31  
    },  
    {  
      "name": "Jill",  
      "age": 29  
    }  
  ]  
}
```

```
0c Buford          name  
72                age (57)  
02                wife: branch 1  
14 Wilhelmina    name  
6e                age (55)  
04                children: element count  
0c Sanjay         name  
3e                age (31)  
08 Jill           name  
3a                age (29)  
00                children: end of array
```

```
0c 42 75 66 6f 72 64 72  
02 14 57 69 6c 68 65 6c  
6d 69 6e 61 6e 04 0c 53  
61 6e 6a 61 79 3e 08 4a  
69 6c 6c 3a 00
```

# Schema resolution

---

- Since writer schema is available, reader schema doesn't have to match perfectly
- Unused record fields are skipped
- Missing record fields must have default values
- Fields, branches, and enum symbols can be reordered
- Promotion of scalar values (`int` → `long`, etc)

# Project information

---

- Language bindings: Java, Python, Ruby, C, C++, etc.
- Website: <http://avro.apache.org/>
- [Specification](#)
- [Mailing lists](#)